ConstantStep MNIST: output-gradient mask comparison

Pulse noise 0 · nominal seed 0 · LR 1.5 · batch 16 · 3/3 runs complete

legacy · n=1/1 native · n=1/1 disabled · n=1/1

All runs complete

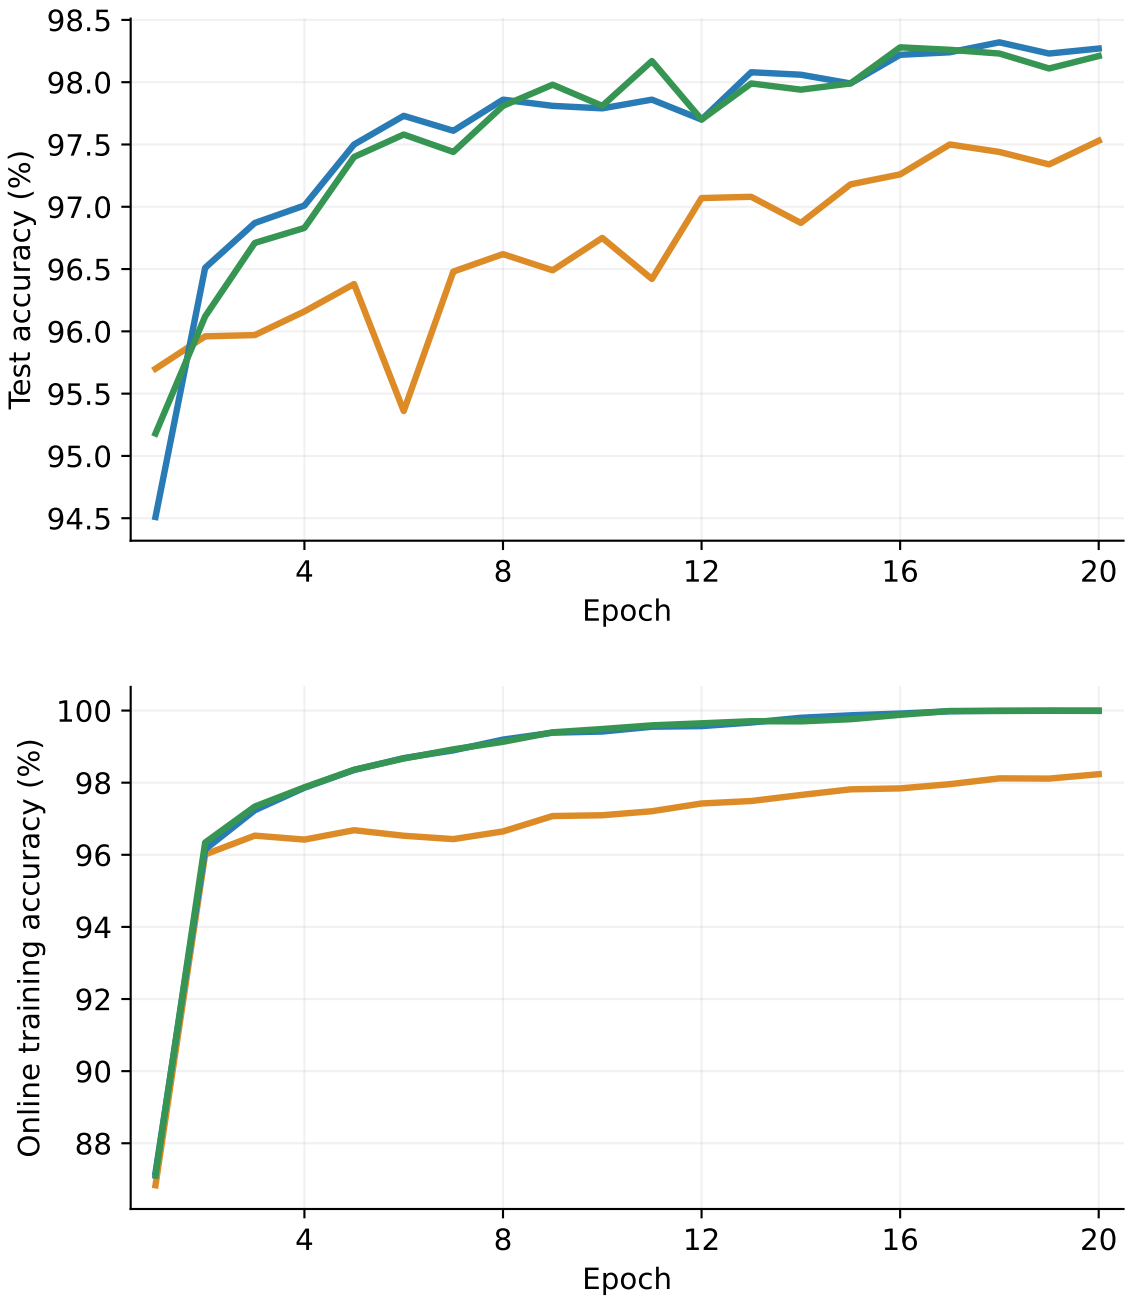

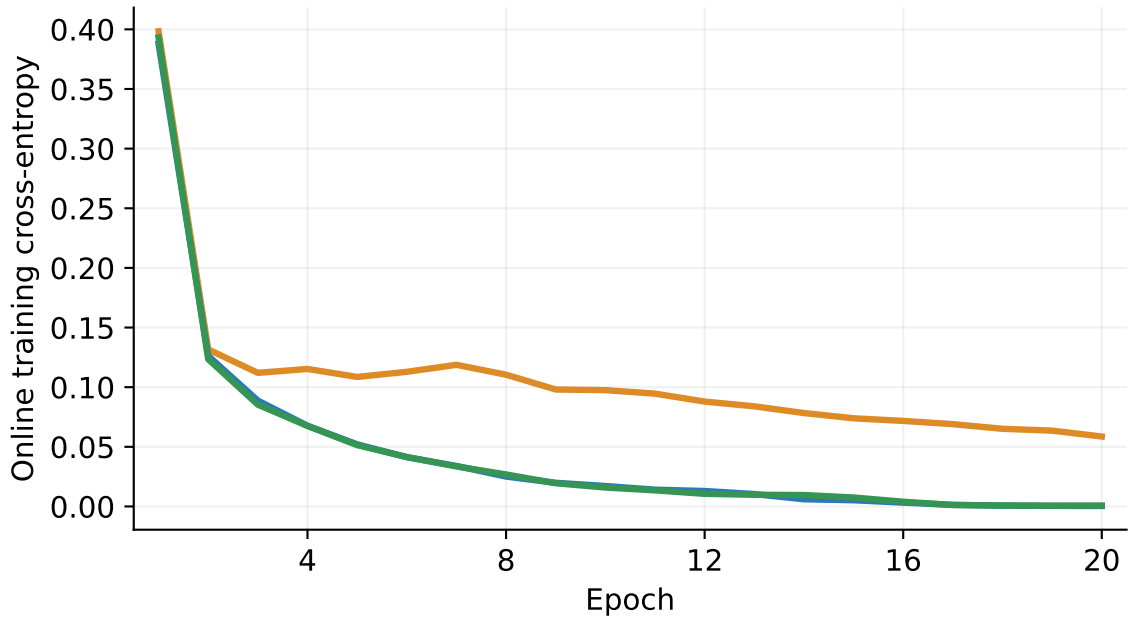

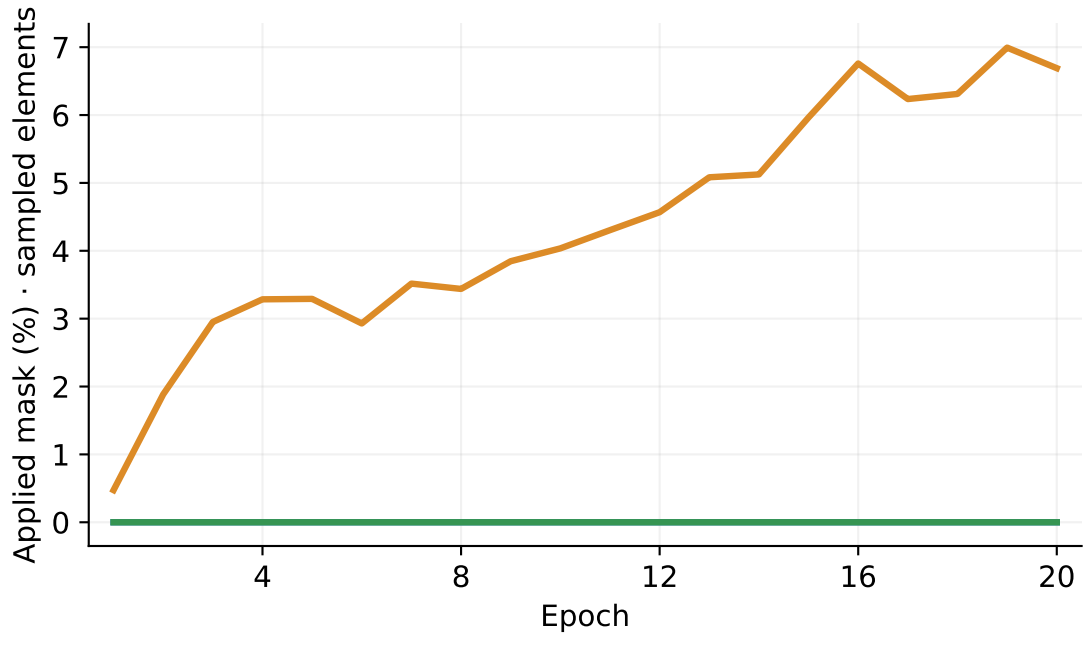

Thick lines: completed-trial mean. Shading: ±1 sample SD (n ≥ 2), not a confidence interval.

Thin lines: individual trials. Dashed: partial runs, excluded from mean/SD. Same seed does not pair backend RNG.

native = this fork's experimental mask. Training metrics are collected during updates; mask rates are sampled.